

Data Mining HW3

20152410 배형준

목차

1. Random Forest Classifier	2
2. Boosting Classifier.....	3
Appendix : R code.....	4
1. Random Forest Classifier	4
2. Boosting Classifier	7

1. Random Forest Classifier

The response is a categorical variable with 10 classes coded from 0 to 9. All predictors are quantitative.

Construct a random forest classifier, report the test classification error and make the confusion matrix. Note “ranger” is faster than “randomForest”.

‘ranger’ 패키지의 ranger 함수를 이용하여 random forest classifier 를 학습하였다. Number of tree 는 default 값인 500 으로 사용해서 학습하였고 tree 를 split 할 때 고려할 변수의 개수인 mtry 도 default 값인 $\sqrt{\text{predictors}} = \sqrt{784} = 28$ 을 사용하였다. 모델 적합 결과 OOB error 는 5.75%, Test classification error 는 6.7%가 나왔다.

Ranger 함수를 cross-validation 과 함께 사용하고 싶다면 ‘spm’ 패키지의 rgcv 함수를 이용할 수 있다. Cross-validation 과 hyper parameter tuning 을 함께 진행하고 싶다면 ‘caret’ 패키지의 train 함수를 trainControl, expand.grid 와 함께 이용하여 모델을 학습시킬 수 있을 것이다.

MNIST 데이터 자체의 power 가 강한 편이고, random forest 모델도 평균 이상의 성능을 보장하는 모델이라 hyper parameter tuning 없이 test error 6.7%를 얻었다. 위에 언급한 방법으로 튜닝을 진행한다면 더 좋은 성능의 모델을 얻을 수 있을 것이다.

2. Boosting Classifier

Construct a boosting classifier, report the test classification error and make the confusion matrix. Note “xgboost” is faster than “gbm”.

‘xgboost’ 패키지의 `xgb.train` 함수를 이용하여 xgboost classifier 모델을 학습시켰다. 더 빠른 학습을 위해 `xgb.DMatrix` 로 데이터를 변환시키는 작업을 거쳤다. Xgboost 의 경우 종속변수가 범주형 변수인 경우에도 integer 로 바꾼 뒤 0 부터 시작하도록 만들어서 `xgb.DMatrix` 에 넣어줘야 한다. 학습 파라미터로 학습률을 0.2, 목적함수로 softmax, loss function 으로 mlogloss 를 이용하였다. 모델이 일정 횟수 이상 base learner 를 추가로 학습했을 때 train 과 test 에 대한 성능이 개선되지 않는다면 모델 학습을 멈추는 `early_stopping_rounds` 를 20 으로 설정해 모델 학습에 불필요한 시간을 단축할 수 있었다.

학습 결과 Test classification error 는 6.5%로 이 데이터에 대해서 random forest classifier 와 비슷한 성능을 보여준다. Rf classifier 에선 class3 f1 score 가 0.8757 로 다른 범주와 비교했을 때 가장 낮았는데, xgb classifier 에선 class3 f1 score 가 0.9048 로 살짝 개선된 모습을 보여준다.

Appendix : R code

1. Random Forest Classifier

Load dataset

```
student = 20152410
```

```
mnist_train = read.csv('./MNIST_train_small.csv', header=TRUE)
```

```
mnist_test = read.csv('./MNIST_test_small.csv', header=TRUE)
```

```
train_data = mnist_train[, 2:785]
```

```
train_label = as.factor(mnist_train$y)
```

```
test_data = mnist_test[, 2:785]
```

```
test_label = as.factor(mnist_test$y)
```

random forest classifier

```
library(caret)
```

```
library(ranger)
```

```
set.seed(student)
```

```
ranger_model = ranger(x = train_data, y = train_label)
```

```
ranger_model
```

```
## Ranger result
```

```
##
```

```
## Call:
```

```
##  ranger(x = train_data, y = train_label)
```

```
##
```

```
## Type: Classification
```

```
## Number of trees: 500
```

```
## Sample size: 6000
```

```
## Number of independent variables: 784
```

```
## Mtry: 28
```

```
## Target node size: 1
```

```
## Variable importance mode: none
```

```
## Splitrule: gini
```

```
## OOB prediction error: 5.75 %
```

```
ranger_pred = predict(ranger_model, data=test_data,  
                      num.trees=ranger_model$num.trees)
```

```
ranger_clf_error = mean(ranger_pred$predictions != test_label)
```

```
cat('Test error of ranger classifier : ', 100*ranger_clf_error, '%')
```

```
## Test error of ranger classifier : 6.7 %
```

```

ranger_table = table(ranger_pred$predictions, test_label)
ranger_cfm = confusionMatrix(ranger_table, mode='everything')
ranger_cfm

## Confusion Matrix and Statistics
##
##      test_label
##      0  1  2  3  4  5  6  7  8  9
## 0  93  0  3  0  0  2  1  0  0  1
## 1  0 106  0  0  0  1  0  0  0  3
## 2  0  0 104  3  0  0  0  2  0  0
## 3  0  0  2  74  0  2  0  0  2  3
## 4  0  0  0  0  99  0  1  1  1  1
## 5  0  0  0  4  0  84  1  0  0  0
## 6  0  0  0  0  1  3  96  0  1  0
## 7  0  0  1  2  0  1  0 100  2  1
## 8  1  0  1  3  0  1  0  1  83  0
## 9  0  0  0  0  8  1  0  4  1  94
##
## Overall Statistics
##
##              Accuracy : 0.933
##              95% CI : (0.9157, 0.9477)
##      No Information Rate : 0.111
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.9255
##
##  Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: 0 Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.9894   1.0000   0.9369   0.8605   0.9167   0.8842
## Specificity      0.9923   0.9955   0.9944   0.9902   0.9955   0.9945
## Pos Pred Value   0.9300   0.9636   0.9541   0.8916   0.9612   0.9438
## Neg Pred Value   0.9989   1.0000   0.9921   0.9869   0.9900   0.9879
## Precision        0.9300   0.9636   0.9541   0.8916   0.9612   0.9438
## Recall           0.9894   1.0000   0.9369   0.8605   0.9167   0.8842
## F1               0.9588   0.9815   0.9455   0.8757   0.9384   0.9130
## Prevalence       0.0940   0.1060   0.1110   0.0860   0.1080   0.0950
## Detection Rate   0.0930   0.1060   0.1040   0.0740   0.0990   0.0840
## Detection Prevalence 0.1000   0.1100   0.1090   0.0830   0.1030   0.0890
## Balanced Accuracy 0.9908   0.9978   0.9657   0.9253   0.9561   0.9393

```

##	Class: 6	Class: 7	Class: 8	Class: 9
## Sensitivity	0.9697	0.9259	0.9222	0.9126
## Specificity	0.9945	0.9922	0.9923	0.9844
## Pos Pred Value	0.9505	0.9346	0.9222	0.8704
## Neg Pred Value	0.9967	0.9910	0.9923	0.9899
## Precision	0.9505	0.9346	0.9222	0.8704
## Recall	0.9697	0.9259	0.9222	0.9126
## F1	0.9600	0.9302	0.9222	0.8910
## Prevalence	0.0990	0.1080	0.0900	0.1030
## Detection Rate	0.0960	0.1000	0.0830	0.0940
## Detection Prevalence	0.1010	0.1070	0.0900	0.1080
## Balanced Accuracy	0.9821	0.9590	0.9573	0.9485

2. Boosting Classifier

boosting classifier

```
library(xgboost)
```

```
train_dmatrix = xgb.DMatrix(data=as.matrix(train_data), label=as.integer(train_label)-1)
```

```
test_dmatrix = xgb.DMatrix(data=as.matrix(test_data), label=as.integer(test_label)-1)
```

```
xgb_params = list(eta=0.2,  
                  num_class=length(levels(train_label)),  
                  objective='multi:softmax',  
                  eval_metric='mlogloss')
```

```
set.seed(student)
```

```
xgb_model = xgb.train(data=train_dmatrix,  
                      params=xgb_params,  
                      nrounds=500,  
                      early_stopping_rounds=20,  
                      watchlist=list(val1=train_dmatrix, val2=test_dmatrix),  
                      verbose=0)
```

```

xgb_model
## ##### xgb.Booster
## raw: 1.5 Mb
## call:
##   xgb.train(params = xgb_params, data = train_dmatrix, nrounds = 500,
##     watchlist = list(val1 = train_dmatrix, val2 = test_dmatrix),
##     verbose = 0, early_stopping_rounds = 20)
## params (as set within xgb.train):
##   eta = "0.2", num_class = "10", objective = "multi:softmax", eval_metric = "mlogloss", silent = "1"
## xgb.attributes:
##   best_iteration, best_msg, best_ntreelimit, best_score, niter
## callbacks:
##   cb.evaluation.log()
##   cb.early.stop(stopping_rounds = early_stopping_rounds, maximize = maximize,
##     verbose = verbose)
## # of features: 784
## niter: 130
## best_iteration : 110
## best_ntreelimit : 110
## best_score : 0.216168
## nfeatures : 784
## evaluation_log:
##   iter val1_mlogloss val2_mlogloss
##     1      1.617869      1.703273
##     2      1.285146      1.410190
## ---
##    129      0.002439      0.216710
##    130      0.002420      0.216638
xgb_pred = as.factor(predict(xgb_model, newdata=test_dmatrix))

xgb_clf_error = mean(xgb_pred != test_label)
cat('Test error of xgboost classifier : ', 100*xgb_clf_error, '%')
## Test error of xgboost classifier : 6.5 %

```



```

xgb_table = table(xgb_pred, test_label)
xgb_cfm = confusionMatrix(xgb_table, mode='everything')
xgb_cfm

## Confusion Matrix and Statistics
##
##           test_label
## xgb_pred  0    1    2    3    4    5    6    7    8    9
##      0  93    0    2    0    0    3    1    1    0    1
##      1    0 106    0    0    0    0    0    0    0    2
##      2    0    0 105    3    0    1    0    1    3    0
##      3    0    0    2  76    0    1    1    0    1    1
##      4    0    0    0    1 101    0    1    0    1    3
##      5    0    0    0    1    0  83    1    0    0    2
##      6    0    0    0    1    0    2  94    0    0    0
##      7    0    0    0    0    0    1    0 100    2    0
##      8    1    0    2    3    0    3    1    1  83    0
##      9    0    0    0    1    7    1    0    5    0  94
##
## Overall Statistics
##
##               Accuracy : 0.935
##               95% CI : (0.9179, 0.9495)
##      No Information Rate : 0.111
##      P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.9277
##
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##               Class: 0 Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.9894   1.0000   0.9459   0.8837   0.9352   0.8737
## Specificity      0.9912   0.9978   0.9910   0.9934   0.9933   0.9956
## Pos Pred Value   0.9208   0.9815   0.9292   0.9268   0.9439   0.9540
## Neg Pred Value   0.9989   1.0000   0.9932   0.9891   0.9922   0.9869
## Precision        0.9208   0.9815   0.9292   0.9268   0.9439   0.9540
## Recall           0.9894   1.0000   0.9459   0.8837   0.9352   0.8737
## F1               0.9538   0.9907   0.9375   0.9048   0.9395   0.9121
## Prevalence       0.0940   0.1060   0.1110   0.0860   0.1080   0.0950
## Detection Rate   0.0930   0.1060   0.1050   0.0760   0.1010   0.0830
## Detection Prevalence 0.1010   0.1080   0.1130   0.0820   0.1070   0.0870
## Balanced Accuracy 0.9903   0.9989   0.9685   0.9386   0.9642   0.9346

```

##	Class: 6	Class: 7	Class: 8	Class: 9
## Sensitivity	0.9495	0.9259	0.9222	0.9126
## Specificity	0.9967	0.9966	0.9879	0.9844
## Pos Pred Value	0.9691	0.9709	0.8830	0.8704
## Neg Pred Value	0.9945	0.9911	0.9923	0.9899
## Precision	0.9691	0.9709	0.8830	0.8704
## Recall	0.9495	0.9259	0.9222	0.9126
## F1	0.9592	0.9479	0.9022	0.8910
## Prevalence	0.0990	0.1080	0.0900	0.1030
## Detection Rate	0.0940	0.1000	0.0830	0.0940
## Detection Prevalence	0.0970	0.1030	0.0940	0.1080
## Balanced Accuracy	0.9731	0.9613	0.9551	0.9485